

Outline

- Definition & main principles
- Several extensions of linear regression
- Trees and forests
- Deep learning

What is it?

- Often considered as a special type of machine learning adapted to
 - Very large datasets
 - Unstructured input data (images, texts, sounds etc.)
- Based on neural network algorithms
- Inputs are replaced by a large number of « features » summarizing the characteristics of the original inputs
- These « features » are automatically generated by the algorithms

Application: recognition of numbers

- 60,000 images of numbers available for training of a predictive algorithm
- 10,000 number images available for testing the trained algorithm

Upload images

```
#Images  
x_train <- mnist$train$x  
#Etiquettes  
y_train <- mnist$train$y
```

Training

```
#Images  
x_test <- mnist$test$x  
x_test_image <- mnist$test$x  
#Labels  
y_test <- mnist$test$y
```

Test

One image of x_train and its label in y_train

```
#Example of number  
plot(as.raster(x_train[26,,]), max=255)  
y_train[26]
```

```
> y_train[26]  
[1] 2
```



Another image x_train and its lable y_train

```
#Exemple chiffre  
plot(as.raster(x_train[62,,]), max=255)  
image(1:28,1:28,NUM)  
y_train[62]
```

```
> y_train[62]  
[1] 4
```



```

> dim(x_train)
[1] 60000    28    28
> x_train[1,,]

```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]	[,11]	[,12]	[,13]	[,14]	[,15]	[,16]	[,17]
[1,]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[2,]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[3,]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[4,]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[5,]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
[6,]	0	0	0	0	0	0	0	0	0	0	0	0	3	18	18	18	126
[7,]	0	0	0	0	0	0	0	0	30	36	94	154	170	253	253	253	253
[8,]	0	0	0	0	0	0	0	49	238	253	253	253	253	253	253	253	253
[9,]	0	0	0	0	0	0	0	18	219	253	253	253	253	253	198	182	247
[10,]	0	0	0	0	0	0	0	0	80	156	107	253	253	205	11	0	43
[11,]	0	0	0	0	0	0	0	0	0	14	1	154	253	90	0	0	0
[12,]	0	0	0	0	0	0	0	0	0	0	0	139	253	190	2	0	0
[13,]	0	0	0	0	0	0	0	0	0	0	0	11	190	253	70	0	0
[14,]	0	0	0	0	0	0	0	0	0	0	0	0	35	241	225	160	108
[15,]	0	0	0	0	0	0	0	0	0	0	0	0	0	81	240	253	253
[16,]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	45	186	253
[17,]	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	16	93

```
# reshape
x_train <- array_reshape(x_train, c(nrow(x_train), 784))
x_test  <- array_reshape(x_test,  c(nrow(x_test), 784))
# rescale
x_train <- x_train / 255
x_test  <- x_test  / 255

y_train <- to_categorical(y_train, 10)
y_test  <- to_categorical(y_test, 10)
```




```
> dim(x_train)
[1] 60000 784
```

	[,129]	[,130]	[,131]	[,132]	[,133]	[,134]	[,135]	[,136]	[,137]	[,138]	[,139]
[1,]	0.0000000	0.0000000	0.0000000	0.0000000	0	0	0	0	0	0	0
[2,]	0.6235294	0.9921569	0.6235294	0.1960784	0	0	0	0	0	0	0
[3,]	0.0000000	0.0000000	0.0000000	0.0000000	0	0	0	0	0	0	0
[4,]	0.0000000	0.0000000	0.0000000	0.0000000	0	0	0	0	0	0	0
[5,]	0.0000000	0.0000000	0.0000000	0.0000000	0	0	0	0	0	0	0
[6,]	0.0000000	0.0000000	0.0000000	0.0000000	0	0	0	0	0	0	0

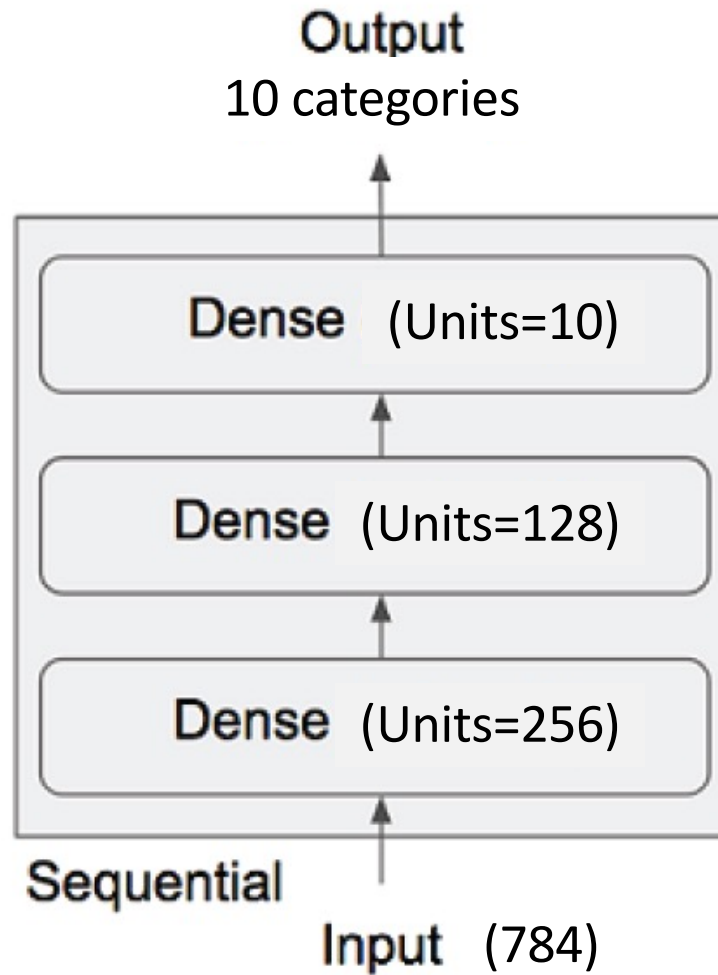
```
> dim(y_train)
[1] 60000 10
```

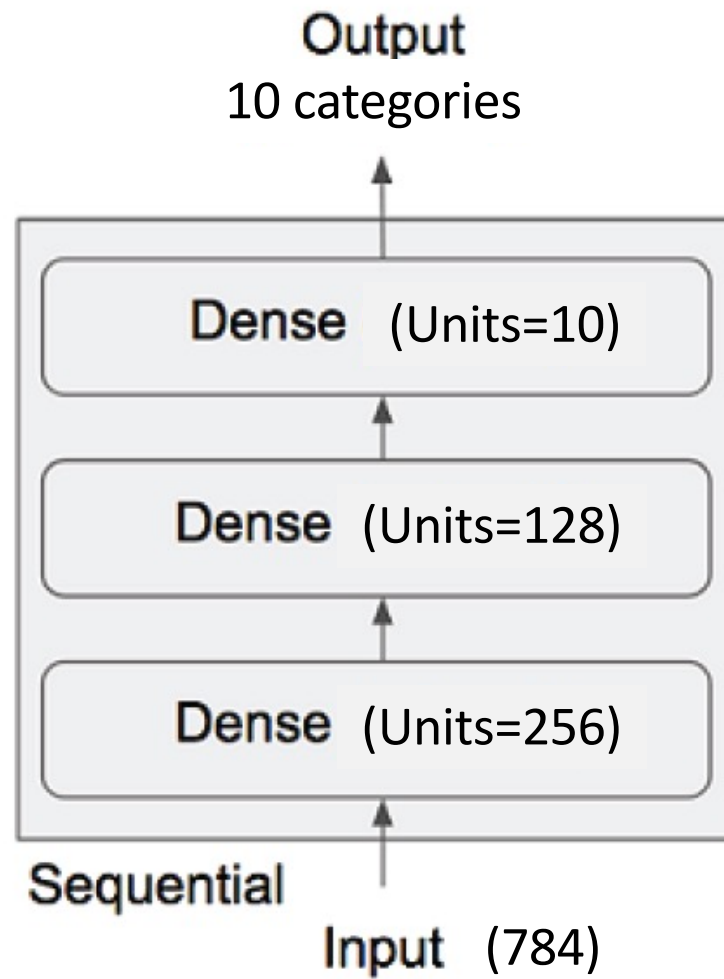
```
> head(y_train)
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]
[1,]	0	0	0	0	0	1	0	0	0	0
[2,]	1	0	0	0	0	0	0	0	0	0
[3,]	0	0	0	0	1	0	0	0	0	0
[4,]	0	1	0	0	0	0	0	0	0	0
[5,]	0	0	0	0	0	0	0	0	0	1
[6,]	0	0	1	0	0	0	0	0	0	0



Neural network with several layers






$$f\left(b3 + \sum_{k=1}^{128} w_{k,3} * X_{k,3}\right) * 10$$

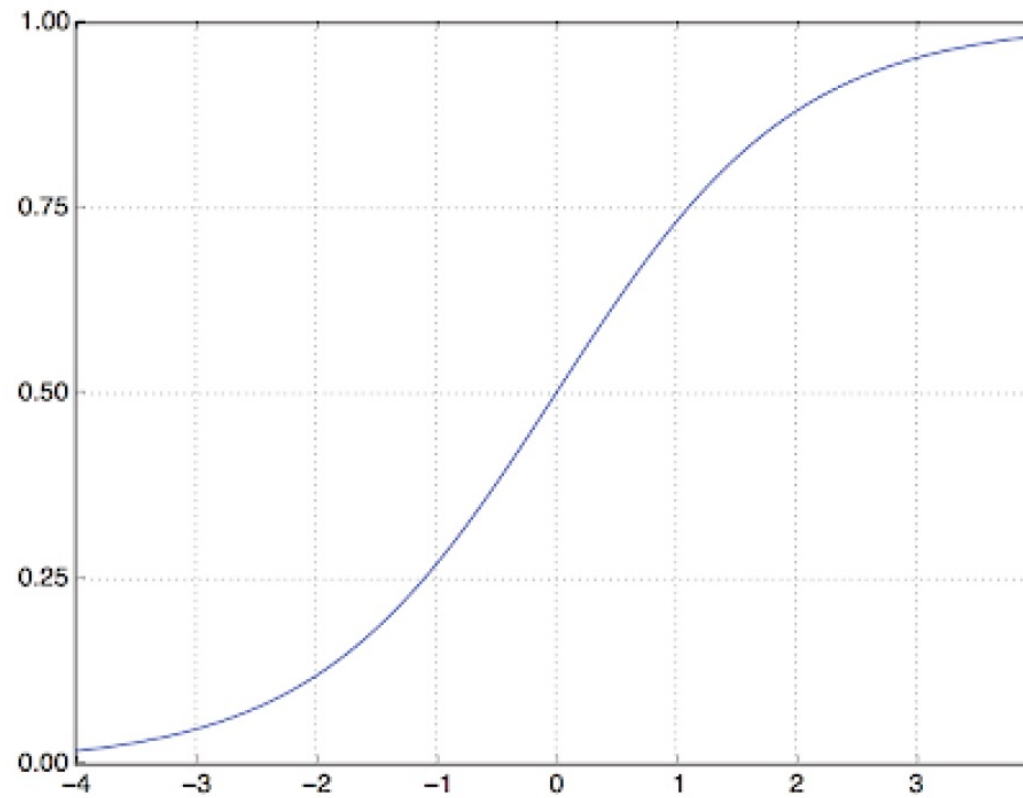
$$f\left(b2 + \sum_{j=1}^{256} w_{j,2} * X_{j,2}\right) * 128$$

$$f\left(b1 + \sum_{i=1}^{784} w_{i,1} * X_{i,1}\right) * 256$$

$f=\text{RELU}$



$f=\text{softmax}$



Neural network with several layers

```
model <- keras_model_sequential()
model %>%

  layer_dense(units = 256, activation = 'relu',
              input_shape = c(784)) %>%

  layer_dropout(rate = 0.4) %>%

  layer_dense(units = 128, activation = 'relu') %>%

  layer_dropout(rate = 0.3) %>%

  layer_dense(units = 10, activation = 'softmax')

summary(model)
```

Many parameters

```
> summary(model)
```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 256)	200960
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 128)	32896
dropout_2 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 10)	1290

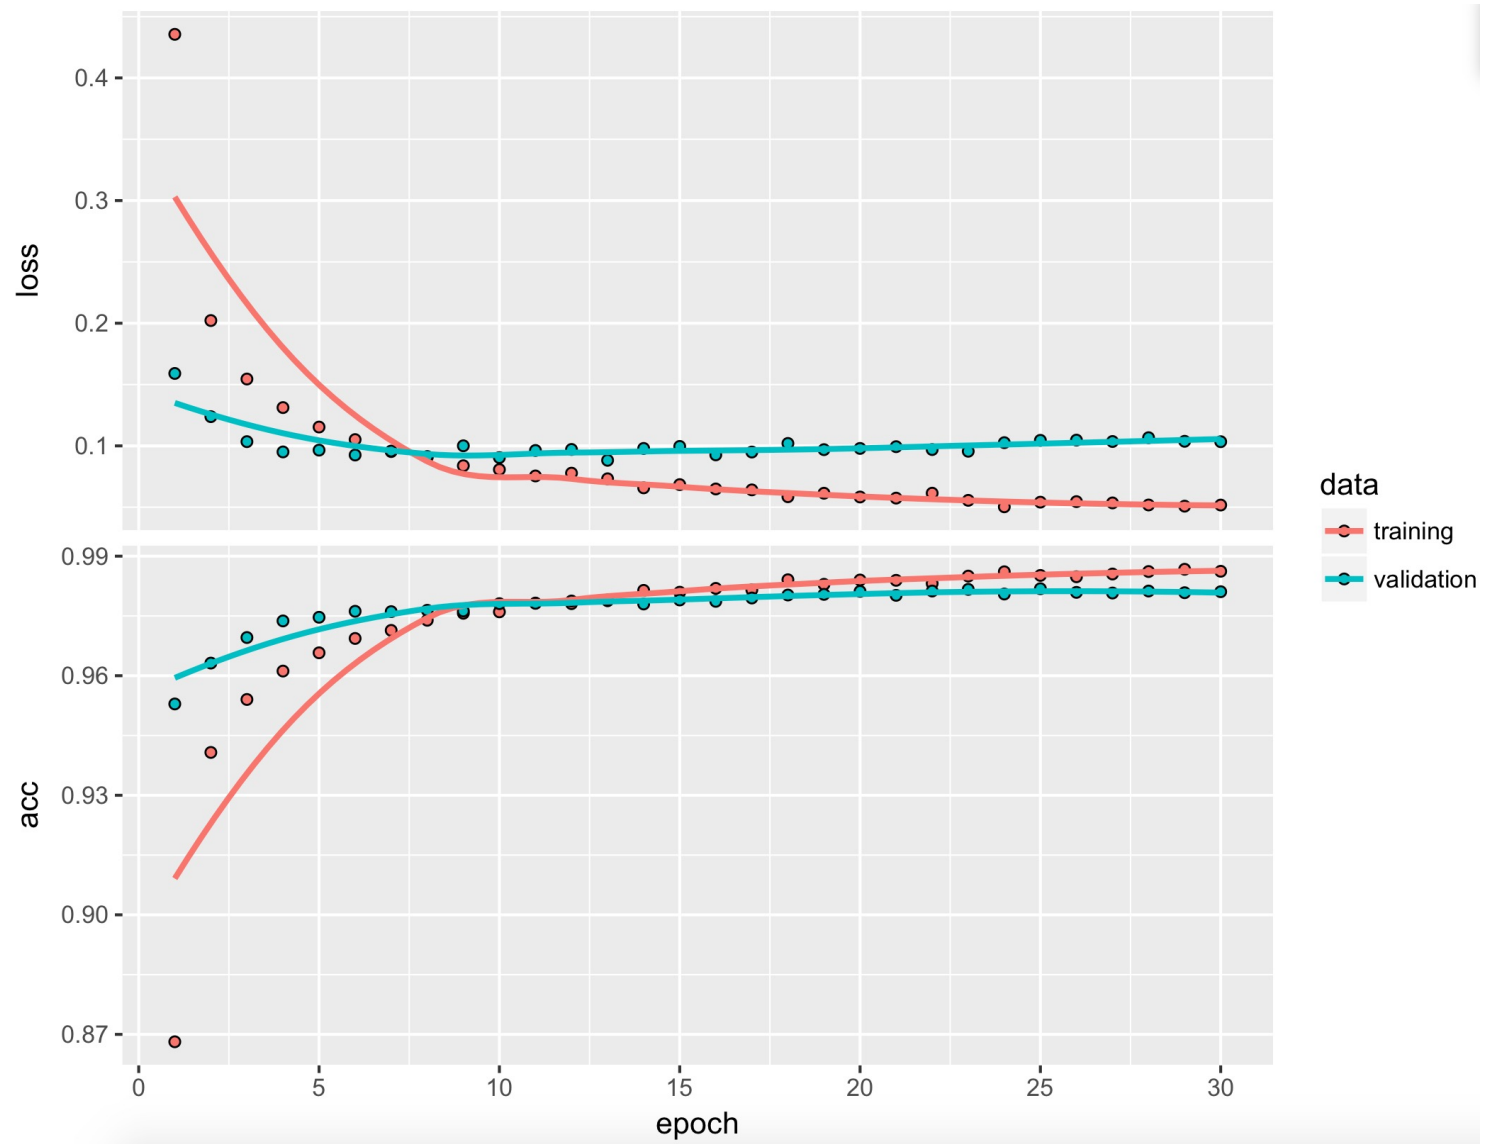
```
Total params: 235,146  
Trainable params: 235,146  
Non-trainable params: 0
```


Model training

```
model %>% compile(  
  loss = 'categorical_crossentropy',  
  optimizer = optimizer_rmsprop(),  
  metrics = c('accuracy')  
)
```

```
history <- model %>% fit(  
  x_train, y_train,  
  epochs = 30, batch_size = 128,  
  validation_split = 0.2  
)
```

```
plot(history)
```



Model testing

```
##Evaluation  
results<- model%>%evaluate(x_test,y_test)  
results
```

```
> results  
$loss  
[1] 0.1017851  
  
$acc  
[1] 0.981
```

```
> Prediction<-model %>% predict_classes(x_test)
> #Exemple
> k<-780
> plot(as.raster(x_test_image[k,,], max=255))
> y_test[k,]
[1] 0 0 0 0 0 1 0 0 0 0
> Prediction[k]
[1] 5
```



Multi-layer neural network with convolutions

What is a convolution?

Convolution=Filter

- Each filter « slides » over the object (e.g. over the pixels of an image) to create a new object called « feature map »
- A feature map shows a specific aspect of an object (e.g., one part of an image)
- Each filter is characterized by a size and by weights
 - ✓ The size is pre-specified,
 - ✓ the weights are optimized like the other model parameters
- The generated feature map is used as a new source of inputs for the neural network
- Several filters are generally used successively, leading to several feature maps

Filter (kernel)

0	1	2
2	2	0
0	1	2

Weight

Dumoulin and Visin 2018. <https://arxiv.org/pdf/1603.07285.pdf>

Filter (kernel)

0	1	2
2	2	0
0	1	2

Weight

3 ₀	3 ₁	2 ₂	1	0
0 ₂	0 ₂	1 ₀	3	1
3 ₀	1 ₁	2 ₂	2	3
2	0	0	2	2
2	0	0	0	1

Original object
(e.g. image)

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

New feature

Filter (kernel)

0	1	2
2	2	0
0	1	2

3_0	3_1	2_2	1	0
0_2	0_2	1_0	3	1
3_0	1_1	2_2	2	3
2	0	0	2	2
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

Filter (kernel)

0	1	2
2	2	0
0	1	2

3_0	3_1	2_2	1	0
0_2	0_2	1_0	3	1
3_0	1_1	2_2	2	3
2	0	0	2	2
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

3	3_0	2_1	1_2	0
0	0_2	1_2	3_0	1
3	1_0	2_1	2_2	3
2	0	0	2	2
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

Filter (kernel)

0	1	2
2	2	0
0	1	2

3_0	3_1	2_2	1	0
0_2	0_2	1_0	3	1
3_0	1_1	2_2	2	3
2	0	0	2	2
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

3	3_0	2_1	1_2	0
0	0_2	1_2	3_0	1
3	1_0	2_1	2_2	3
2	0	0	2	2
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

3	3	2_0	1_1	0_2
0	0	1_2	3_2	1_0
3	1	2_0	2_1	3_2
2	0	0	2	2
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

Dumoulin and Visin 2018. <https://arxiv.org/pdf/1603.07285.pdf>

Filter (kernel)

0	1	2
2	2	0
0	1	2

3_0	3_1	2_2	1	0
0_2	0_2	1_0	3	1
3_0	1_1	2_2	2	3
2	0	0	2	2
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

3	3_0	2_1	1_2	0
0	0_2	1_2	3_0	1
3	1_0	2_1	2_2	3
2	0	0	2	2
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

3	3	2_0	1_1	0_2
0	0	1_2	3_2	1_0
3	1	2_0	2_1	3_2
2	0	0	2	2
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

1.7	1.7	1.7
1.0	1.2	1.8
1.1	0.8	1.3

Dumoulin and Visin 2018. <https://arxiv.org/pdf/1603.07285.pdf>

```
x_train <- mnist$train$x  
y_train <- mnist$train$y  
x_test  <- mnist$test$x  
y_test  <- mnist$test$y
```

```
# rescale  
x_train <- x_train / 255  
x_test  <- x_test  / 255
```

```
x_train<-array(x_train,dim=c(60000,28,28,1))  
x_test<-array(x_test,dim=c(10000,28,28,1))
```

```
y_train <- to_categorical(y_train, 10)  
y_test  <- to_categorical(y_test, 10)
```

```
model <- keras_model_sequential()
model %>%
  layer_conv_2d(filters = 32, kernel_size = c(3,3), activation = 'relu',
                input_shape = c(28,28,1)) %>%
  layer_conv_2d(filters = 32, kernel_size = c(3,3), activation = 'relu') %>%
  layer_max_pooling_2d(pool_size = c(2,2)) %>%
  layer_dropout(rate = 0.25) %>%
  layer_conv_2d(filters = 64, kernel_size = c(3,3), activation = 'relu') %>%
  layer_conv_2d(filters = 64, kernel_size = c(3,3), activation = 'relu') %>%
  layer_max_pooling_2d(pool_size = c(2,2)) %>%
  layer_dropout(rate = 0.25) %>%
  layer_flatten() %>%
  layer_dense(units = 256, activation = 'relu') %>%
  layer_dropout(rate = 0.25) %>%
  layer_dense(units = 10, activation = 'softmax')

summary(model)
```

```
model %>% compile(  
  loss = 'categorical_crossentropy',  
  optimizer = optimizer_rmsprop(),  
  metrics = c('accuracy')  
)
```

```
history <- model %>% fit(  
  x_train, y_train,  
  epochs = 2, batch_size = 128,  
  validation_split = 0.2  
)
```

```
plot(history)
```

```
> Prediction<-model %>% predict_classes(x_test)
> #Exemple
> k<-780
> plot(as.raster(x_test_image[k,,], max=255))
> y_test[k,]
[1] 0 0 0 0 0 1 0 0 0 0
> Prediction[k]
[1] 5
```



More and more applications

- Classification of objects (numbers, plants, animals, words etc.)
- Diagnosis of diseases
- Weather forecasts
- Automatic translation of text

ARTICLE OPEN

Deep learning enables robust assessment and selection of human blastocysts after in vitro fertilization

Pegah Khosravi^{1,2}, Ehsan Kazemi³, Qiansheng Zhan⁴, Jonas E. Malmsten⁴, Marco Toschi⁴, Pantelis Zisimopoulos^{1,2}, Alexandros Sigaras^{1,2}, Stuart Lavery⁵, Lee A. D. Cooper⁶, Cristina Hickman⁵, Marcos Meseguer⁷, Zev Rosenwaks⁴, Olivier Elemento^{1,2,8}, Nikica Zaninovic⁴ and Iman Hajirasouliha^{1,2}




Fig. 4 STORK vs. embryologists classification: STORK classifies the fair-quality images into existing good-quality and poor-quality classes. For example, panels "a" and "b" are labeled 3A-B (fair-quality) according to the Veeck and Zaninovic grading system, while STORK classified them as poor-quality and good-quality, respectively. Also, panels "c" and "d" are both labeled 3BB (fair-quality). However, the algorithm correctly classified panel "c" as poor-quality and panel "d" as good-quality. As the figure shows, the outcome in the embryos in "b" and "d" is positive live birth, whereas it is negative live birth in "a" and "c"

Outline

- Brief overview of machine learning
- Trees and forests
- Deep learning
- **Conclusion**

Main challenges in machine learning projects

- Choose a relevant question (Which Y? Which X?)
- Find reliable data
- Calibrate the hyper-parameters
- Assess prediction accuracy without bias
- Optimize computation time
- Vizualisation of output responses

Start simple

Start with two simple methods:

- Penalized linear regression (ex: LASSO)
- Random forest

Some trends

- Visualization tools (to open « the black boxes »)
- Image and text analyses (text mining, deep learning)
- Packages to streamline the development of predictive models (keras, caret, H2O...)
- Including expert knowledge in machine learning

Machine learning to emulate complexe models



ELSEVIER

Contents lists available at [ScienceDirect](#)

Journal of Environmental Management

journal homepage: www.elsevier.com/locate/jenvman

Research article

Meta-modeling methods for estimating ammonia volatilization from nitrogen fertilizer and manure applications

Maharavo Marie Julie Ramanantenaso^{a,b}, Sophie Générumont^{a,*}, Jean-Marc Gilliot^a, Carole Bedos^a, David Makowski^{c,d}

ACCEPTED MANUSCRIPT • OPEN ACCESS

Maize yield and nitrate loss prediction with machine learning algorithms

Mohsen Shahhosseini¹, Rafael A Martinez-Feria² , Guiping Hu³ and Sotirios Archontoulis¹

Accepted Manuscript online 29 October 2019 • © 2019 The Author(s). Published by IOP Publishing Ltd